

CS 444/544 Compilers
Ozyegin University
<http://aktemur.github.io/cs544>

Note by Baris Aktemur:

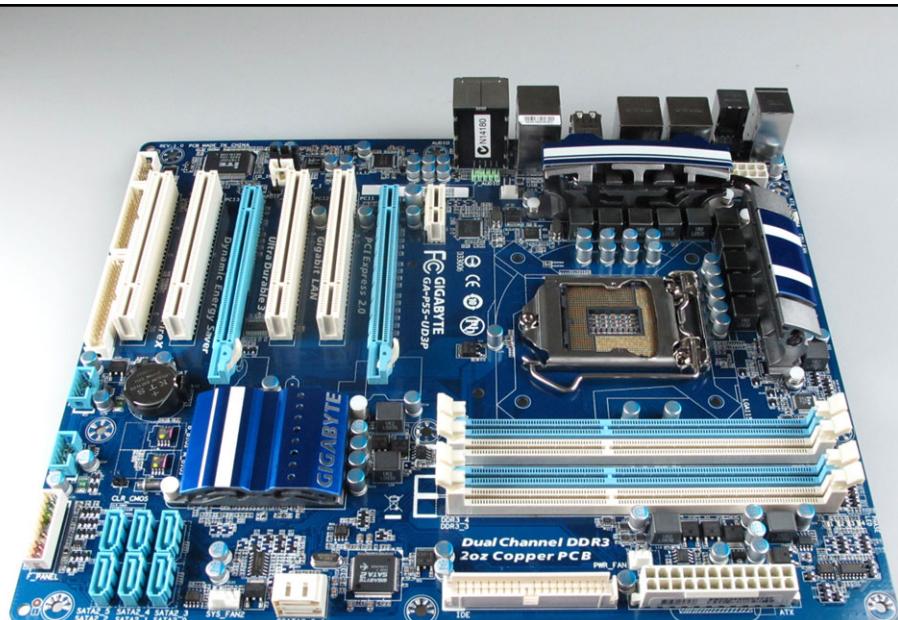
Our slides are adapted from Cooper and Torczon's slides that they prepared for COMP 412 at Rice.

Copyright 2010, Keith D. Cooper & Linda Torczon, all rights reserved.

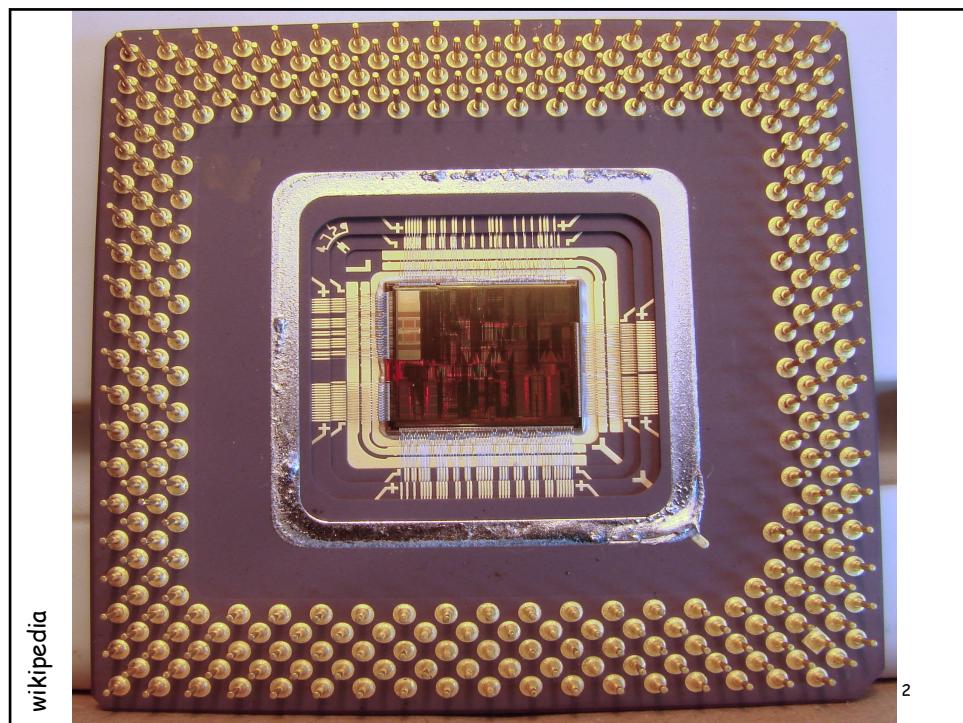
Students enrolled in Comp 412 at Rice University have explicit permission to make copies of these materials for their personal use.

Faculty from other educational institutions may use these materials for nonprofit educational purposes, provided this copyright notice is preserved.

0

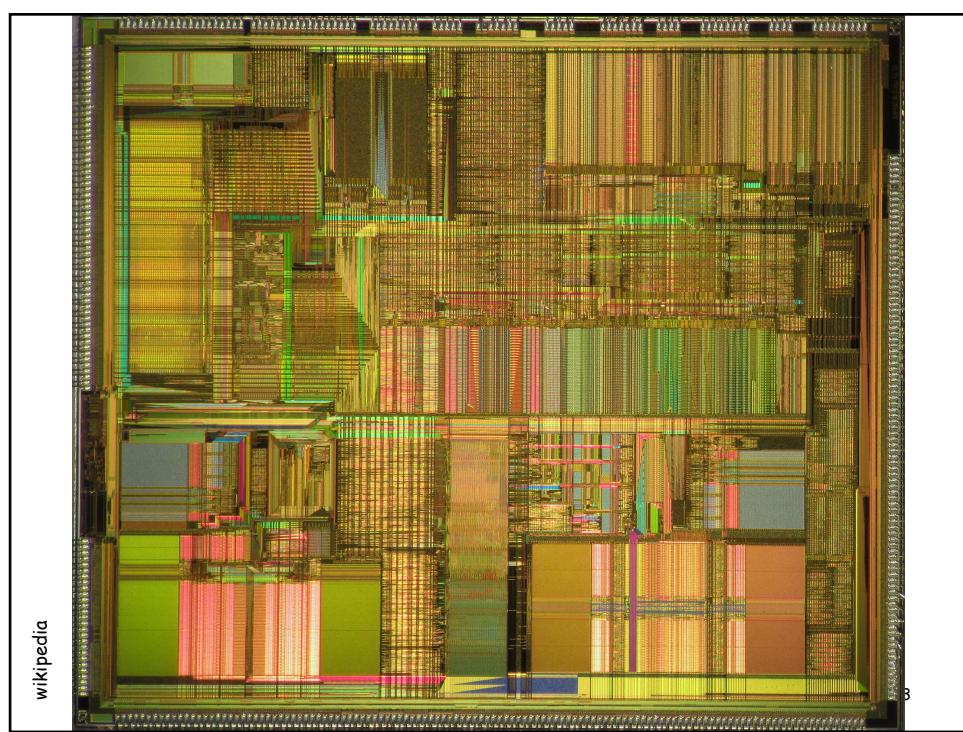


nextmedia.com.au



wikipedia

2



wikipedia

3

Instructing a CPU

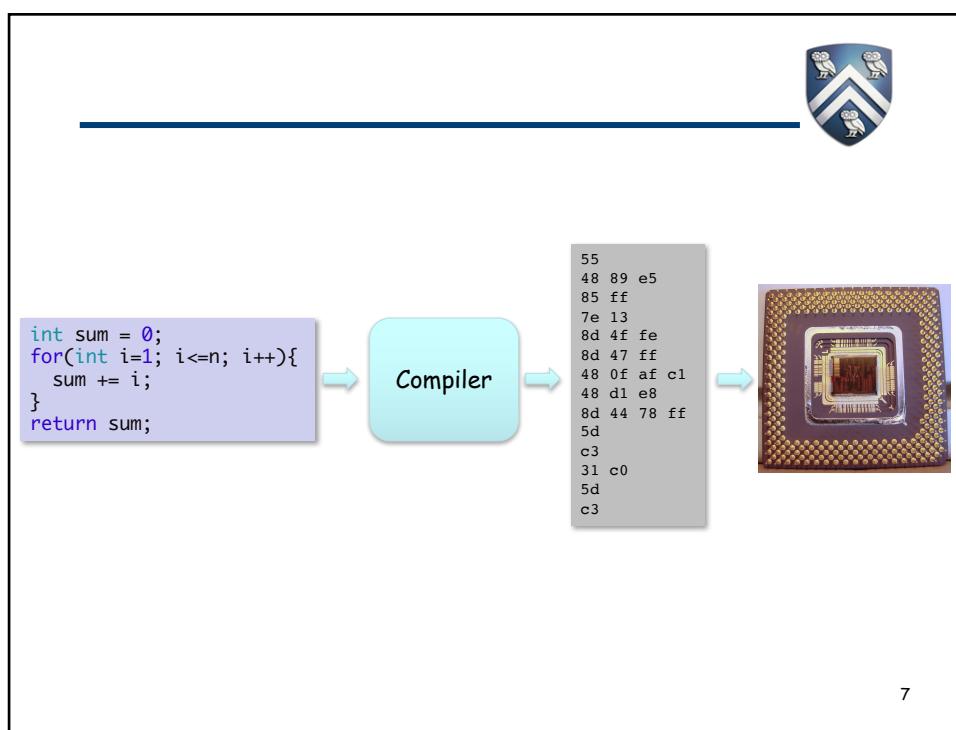
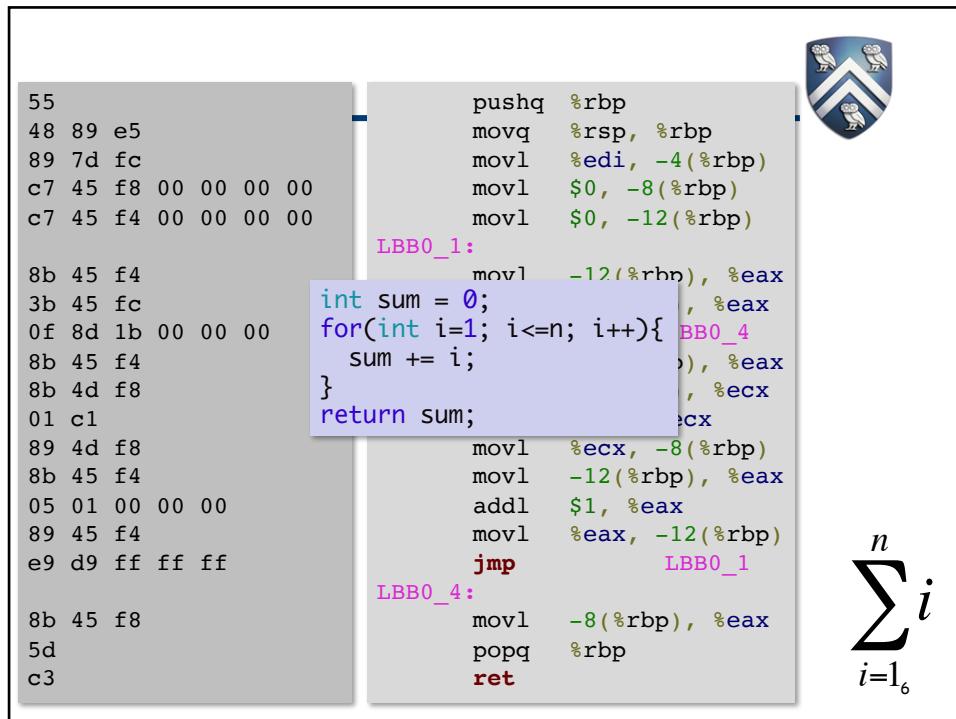


- Wires go in, wires come out
 - Low voltage, high voltage
 - 0 and 1

$$\sum_{i=1}^n i$$

4

$$\sum_{i=1_5}^n i$$



Compilers



- What is a **compiler**?
 - A program that translates an *executable* program in one language into an *executable* program in another language
 - The compiler should improve the program, *in some way*
- What is an **interpreter**?
 - A program that reads an *executable* program and produces the results of executing that program
- C,C++ are typically compiled. Scheme, Python are typically interpreted. Haskell, F#, Ocaml are compiled or interpreted, depending on the context.
- Java is compiled to bytecodes (code for the Java VM)
 - which are then interpreted
 - Or a hybrid strategy is used
 - Just-in-time compilation

Common mis-statement:
*X is an interpreted language
 (or a compiled language)*

Comp 412, Fall 2010

8

Why Study Compilation?



- Compilers are **important**
 - Responsible for many aspects of system performance
 - Attaining performance has become more difficult over time
 - In 1980, typical code got 85% or more of peak performance
 - Today, that number is closer to 5 to 10% of peak
 - Compiler has become a prime determiner of performance
- Compilers are **interesting**
 - Compilers include many applications of theory to practice
 - Writing a compiler exposes algorithmic & engineering issues
- Compilers are **everywhere**
 - Many practical applications have embedded languages
 - Commands, macros, formatting tags ...
 - Many applications have input formats that look like languages

Comp 412, Fall 2010

9

Reducing the Price of Abstraction



Computer Science is the art of creating virtual objects and making them useful.

- We invent abstractions and uses for them
- We invent ways to make them efficient
- Programming is the way we realize these inventions

Well written compilers make abstraction affordable

- Cost of executing code should reflect the underlying work rather than the way the programmer chose to write it
- Change in expression should bring small performance change
- Cannot expect compiler to devise better algorithms
 - Don't expect bubblesort to become quicksort

Comp 412, Fall 2010

10

Making Languages Usable



It was our belief that if FORTRAN, during its first months, were to translate any reasonable "scientific" source program into an object program only half as fast as its hand-coded counterpart, then acceptance of our system would be in serious danger... I believe that had we failed to produce efficient programs, the widespread use of languages like FORTRAN would have been seriously delayed.

— John Backus on the subject of the 1st FORTRAN compiler

Comp 412, Fall 2010

11

Simple Examples

All data collected with gcc 4.1, -O3, running on a quiescent, multiuser Intel T9600 @ 2.8 GHz



Which is faster?

All three loops have distinct performance.

0.51 sec on 10,000 × 10,000 array

```
for (i=0; i<n; i++)
    for (j=0; j<n; j++)
        A[i][j] = 0;
```

1.65 sec on 10,000 × 10,000 array

```
for (i=0; i<n; i++)
    for (j=0; j<n; j++)
        A[j][i] = 0;
```

0.11 sec on 10,000 × 10,000 array

```
p = &A[0][0];
t = n * n;
for (i=0; i<t; i++)
    *p++ = 0;
```

A good compiler should know these tradeoffs, on each target, and generate the best code.
Few real compilers do.

Conventional wisdom suggests using

```
bzero((void*) &A[0][0],(size_t) n*n*sizeof(int))
```

0.52 sec on 10,000 × 10,000 array

Comp 412, Fall 2010 12

Simple Examples

Example from Rⁿ Programming Environment, Rice, circa 1984



Abstraction has its price (& it is often higher than expected)

```
struct point { /* Point on the plane of windows */
    int x; int y;
}

void Padd(struct point p, struct point q, struct point * r)
{
    r->x = p.x + q.x;
    r->y = p.y + q.y;
}

int main( int argc, char *argv[] )
{
    struct point p1, p2, p3;

    p1.x = 1; p1.y = 1;
    p2.x = 2; p2.y = 2;

    Padd(p1, p2, &p3);

    printf("Result is <%d,%d>.n", p3.x, p3.y);
}
```

Comp 412, Fall 2010 13

Simple Example (point add)

gcc 4.1, -S option



```

main: (some boilerplate code elided for brevity's sake)
L5:
    popl %ebp
    movl $1, -16(%ebp)
    movl $1, -12(%ebp)
    movl $2, -24(%ebp)
    movl $2, -20(%ebp)
    leal -32(%ebp), %eax
    movl %eax, 16(%esp)
    movl -24(%ebp), %eax
    movl -20(%ebp), %edx
    movl %eax, 8(%esp)
    movl %edx, 12(%esp)
    movl -16(%ebp), %eax
    movl -12(%ebp), %edx
    movl %eax, (%esp)
    movl %edx, 4(%esp)
    call PAdd
    movl -28(%ebp), %eax
    movl -32(%ebp), %edx
    movl %eax, 8(%esp)
    movl %edx, 4(%esp)
    leal L0-"L000000000001$pb"(%ebx), %eax
    movl %eax, (%esp)
    call L_printf$stub
    addl $68, %esp
    popl %ebp
    leave
    ret

```

Code for Intel Core 2 Duo

Assignments to p1 and p2

Setup for call to PAdd

Setup for call to printf

Address calculation for format string in printf call

14

Simple Example (point add)

gcc 4.1, -S option



```

_PAdd:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl 8(%ebp), %edx
    movl 16(%ebp), %eax
    addl %eax, %edx
    movl 24(%ebp), %eax
    movl %edx, (%eax)
    movl 12(%ebp), %edx
    movl 20(%ebp), %eax
    addl %eax, %edx
    movl 24(%ebp), %eax
    movl %edx, 4(%eax)
    leave
    ret

```

Code for PAdd

Actual work

The code does a lot of work to execute two add instructions.
→ Factor of 10 in overhead
→ And a window system does a lot of point adds
Code optimization (careful compile-time reasoning & transformation) can make matters better.

Comp 412, Fall 2010

15

Simple Example (point add)

gcc 4.1, -S -O3 option



```
main: (some boilerplate code ellided for brevity's sake)
L5:
    popl %ebx
    subl $20,%esp
    movl $3,8(%esp)
    movl $3,4(%esp)
    leal LC0-"L00000000001$pb"(%ebx),%eax
    movl %eax,(%esp)
    call L_printf$stub
    addl $20,%esp
    popl %ebx
    leave
    ret
```

It inlined PAdd and folded the known constant values of p1 and p2.

With the right information, a good compiler can work wonders.

→ It kept the body of PAdd because it could not tell if it was dead

What if it could not discern the values of p1 and p2?

Comp 412, Fall 2010

16

Simple Example (point add)

gcc 4.1, -S -O3 option



```
main: (some boilerplate code ellided for brevity's sake)
L5:
    popl %ebx
    subl $20,%esp
    movl _one-"L00000000001$pb"(%ebx),%eax
    addl _two-"L00000000001$pb"(%ebx),%eax
    movl %eax,8(%esp)
    movl %eax,4(%esp)
    leal LC0-"L00000000001$pb"(%ebx),%eax
    movl %eax,(%esp)
    call L_printf$stub
    addl $20,%esp
    popl %ebx
    leave
    ret
```

I put 1 and 2 in global variables named "one" and "two".

The optimizer inlined PAdd

The optimizer recognized that
 $p1.x = p1.y$ and $p2.x = p2.y$
 so

$$p1.x + p2.x = p1.y + p2.y.$$

If I make PAdd static, it
 deletes the code for PAdd

This code shows the more general version. It inlined PAdd and subjected the arguments to local optimization. It still had to perform the adds, but it recognized that the second one was redundant.

→ Gcc did a good job on this example.

Comp 412, Fall 2010

17

Intrinsic Merit



- Compiler construction poses challenging and interesting problems:
 - Compilers must process large inputs, perform complex algorithms, but also **run quickly**
 - Compilers have primary responsibility for **run-time performance**
 - Compilers are responsible for making it acceptable to use the **full power** of the programming language
 - Computer architects perpetually create new challenges for the compiler by building more **complex machines**
 - Compilers must hide that complexity from the programmer
- A successful compiler requires mastery of the many complex interactions between its constituent parts

Comp 412, Fall 2010

18

Intrinsic Interest



- Compiler construction involves ideas from many different parts of computer science

<i>Artificial intelligence</i>	Greedy algorithms Heuristic search techniques
<i>Algorithms</i>	Graph algorithms, union-find Dynamic programming
<i>Theory</i>	DFAs & PDAs, pattern matching Fixed-point algorithms
<i>Systems</i>	Allocation & naming, Synchronization, locality
<i>Architecture</i>	Pipeline & hierarchy management Instruction set use

Comp 412, Fall 2010

19

Why Does This Matter Today?



In the last years, most processors have gone multicore

- The era of clock-speed improvements is drawing to an end
 - Faster clock speeds mean higher power (n^2 effect)
 - Smaller wires mean higher resistance for on-chip wires
- For the near term, performance improvement will come from placing multiple copies of the processor (*core*) on a single die
 - Classic programs, written in old languages, are not well suited to capitalize on this kind of multiprocessor parallelism
 - Parallel languages, some kinds of OO systems, functional languages
 - Parallel programs require sophisticated compilers
- Think of the Intel/AMD bet on multicore as a full-employment act for well-trained compiler writers